

Class 19: Investigating Pertussis Resurgence

Mini Project

Mankeerat Rataul

Import the case data for pertussis below:

Investigating Pertussis cases by year

Q1. With the help of the R “addin” package datapasta assign the CDC pertussis case number data to a data frame called cdc and use ggplot to make a plot of cases numbers over time.

```
cdc <- read.csv("U.S. Reported Pertussis Cases_ 1922 - 2025.csv")
head(cdc)
```

	Year	Number.of.Reported.Pertussis.Cases	Data.Status
1	1922	107,473	Finalized
2	1923	164,191	Finalized
3	1924	165,418	Finalized
4	1925	152,003	Finalized
5	1926	202,210	Finalized
6	1927	181,411	Finalized


There is a problem where half of the data is in character form so we need to transform the Reported Case column to integer form.

```
cdc$Number.of.Reported.Pertussis.Cases <- as.integer(gsub(",", "", cdc$Number.of.Reported.Pertussis.Cases))
head(cdc)
```

	Year	Number.of.Reported.Pertussis.Cases	Data.Status
1	1922	107473	Finalized
2	1923	164191	Finalized
3	1924	165418	Finalized
4	1925	152003	Finalized
5	1926	202210	Finalized
6	1927	181411	Finalized

```
library(ggplot2)

ggplot(cdc, aes(x=Year, y=Number.of.Reported.Pertussis.Cases)) +
  geom_point() +
  geom_line() +
  labs(title = "Pertussis Cases in the US (1922-2025)",
       x = "Year",
       y = "Number of Cases") +
  theme_minimal()
```




A tale of two vaccines (wp and ap)

Q2. Using the ggplot `geom_vline()` function add lines to your previous plot for the 1946 introduction of the wP vaccine and the 1996 switch to aP vaccine (see

example in the hint below). What do you notice?

```
library(ggplot2)

ggplot(cdc, aes(x=Year, y=Number.of.Reported.Pertussis.Cases)) +
  geom_point() +
  geom_line() +
  geom_vline(xintercept = 1946, linetype = "dashed", color = "red") +
  geom_vline(xintercept = 1996, linetype = "dashed", color = "blue") +
  labs(title = "Pertussis Cases in the US (1922-2025)",
       x = "Year",
       y = "Number of Cases") +
  theme_minimal()
```



Q3. Describe what happened after the introduction of the aP vaccine? Do you have a possible explanation for the observed trend?

It seems like cases started spiking after the introduction. It could be that the vaccine is possibly less effective, however it could also be due to other factors like vaccine hesitancy as mentioned in class.

Exploring CMI-PB data

Jsonlite needs to be called to read the CMI-PB data below.

```
library(jsonlite)

subject <- read_json("https://www.cmi-pb.org/api/subject", simplifyVector=TRUE)

head(subject)
```

	subject_id	infancy_vac	biological_sex	ethnicity	race
1	1	wP	Female	Not Hispanic or Latino	White
2	2	wP	Female	Not Hispanic or Latino	White
3	3	wP	Female	Unknown	White
4	4	wP	Male	Not Hispanic or Latino	Asian
5	5	wP	Male	Not Hispanic or Latino	Asian
6	6	wP	Female	Not Hispanic or Latino	White

	year_of_birth	date_of_boost	dataset
1	1986-01-01	2016-09-12	2020_dataset
2	1968-01-01	2019-01-28	2020_dataset
3	1983-01-01	2016-10-10	2020_dataset
4	1988-01-01	2016-08-29	2020_dataset
5	1991-01-01	2016-08-29	2020_dataset
6	1988-01-01	2016-10-10	2020_dataset

Q4. How many aP and wP infancy vaccinated subjects are in the dataset?

```
table(subject$infancy_vac)
```

```
aP wP
87 85
```

87 aP subjects and 85 wP subjects are in the dataset, for a total of 172 subjects.

Q5. How many Male and Female subjects/patients are in the dataset?

```
table(subject$biological_sex)
```

```
Female  Male
112     60
```

112 female and 60 male subjects.

Q6. What is the breakdown of race and biological sex (e.g. number of Asian females, White males etc...)?

```
table(subject$biological_sex, subject$race)
```

	American Indian/Alaska Native	Asian	Black or African American
Female	0	32	2
Male	1	12	3

	More Than One Race	Native Hawaiian or Other Pacific Islander
Female	15	1
Male	4	1

	Unknown or Not Reported	White
Female	14	48
Male	7	32

32 asian females, 32 white males, etc. for other races

Below I'll use a library to deal with dates:

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

date, intersect, setdiff, union

```
today()
```

```
[1] "2026-05-28"
```

Q7. Using this approach determine (i) the average age of wP individuals, (ii) the average age of aP individuals; and (iii) are they significantly different?

So we can use `ymd()` to convert the date of birth from a string into a date and then subtract from the current date. I'm gonna need to use tidyverse in order to filter the pipe for each respective category

```
subject$age <- time_length(today() - ymd(subject$year_of_birth), "years")

library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr 1.2.1      v stringr 1.6.0
v forcats 1.0.1    v tibble 3.3.1
v purrr 1.2.2      v tidyr 1.3.2
v readr 2.2.0

-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x purrr::flatten() masks jsonlite::flatten()
x dplyr::lag() masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
##First gonna find the one for ap vaccine
apdataset <- subject %>%
  filter(infancy_vac == "aP")

wpdataset <- subject %>%
  filter(infancy_vac == "wP")

mean(apdataset$age)
```

```
[1] 28.29918
```

```
mean(wpdataset$age)
```

```
[1] 37.04989
```

```
ttest_result <- t.test(apdataset$age, wpdataset$age)

ttest_result
```

Welch Two Sample t-test

```
data: apdataset$age and wpdataset$age
t = -12.918, df = 104.03, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -10.094058 -7.407351
sample estimates:
mean of x mean of y
 28.29918  37.04989
```

The average age of subjects in the ap dataset is ~28.30 years old, and for the wp dataset it is 37.05 years old.

According to the t-test, the p value is extremely low thus it's a clear difference in age between these groups.

Q8. Determine the age of all individuals at time of boost?

```
subject$age_at_boost <- time_length(ymd(subject$date_of_boost) - ymd(subject$year_of_birth),
head(subject)
```


	subject_id	infancy_vac	biological_sex	ethnicity	race
1	1	wP	Female	Not Hispanic or Latino	White
2	2	wP	Female	Not Hispanic or Latino	White
3	3	wP	Female	Unknown	White
4	4	wP	Male	Not Hispanic or Latino	Asian
5	5	wP	Male	Not Hispanic or Latino	Asian
6	6	wP	Female	Not Hispanic or Latino	White

	year_of_birth	date_of_boost	dataset	age	age_at_boost
1	1986-01-01	2016-09-12	2020_dataset	40.40246	30.69678
2	1968-01-01	2019-01-28	2020_dataset	58.40383	51.07461
3	1983-01-01	2016-10-10	2020_dataset	43.40315	33.77413
4	1988-01-01	2016-08-29	2020_dataset	38.40383	28.65982
5	1991-01-01	2016-08-29	2020_dataset	35.40315	25.65914
6	1988-01-01	2016-10-10	2020_dataset	38.40383	28.77481

Q9. With the help of a faceted boxplot or histogram (see below), do you think these two groups are significantly different?

```
ggplot(subject) +
  aes(age, fill = infancy_vac) +
  geom_histogram(show.legend=FALSE) +
  facet_wrap(vars(infancy_vac), nrow=2) +
  xlab("Age in years")
```

`stat\_bin()` using `bins = 30`. Pick better value `binwidth`.



It does seem like there is a big difference in these groups by histogram with wP being administered to those generally older whereas aP is rarely administered to anyone older than 30.

Below I need to download and append the other data to this dataset in order to interpret antibody titers, etc.

```
specimen <- read_json("https://www.cmi-pb.org/api/specimen", simplifyVector = TRUE)
titer <- read_json("https://www.cmi-pb.org/api/plasma_ab_titer", simplifyVector = TRUE)
```

Q9. Complete the code to join specimen and subject tables to make a new merged data frame containing all specimen records along with their associated subject details:


```
head(specimen)
```

	specimen_id	subject_id	actual_day_relative_to_boost
1	1	1	-3
2	2	1	1
3	3	1	3
4	4	1	7
5	5	1	11
6	6	1	32

	planned_day_relative_to_boost	specimen_type	visit
1	0	Blood	1
2	1	Blood	2
3	3	Blood	3
4	7	Blood	4
5	14	Blood	5
6	30	Blood	6

```
head(subject)
```

	subject_id	infancy_vac	biological_sex	ethnicity	race
1	1	wP	Female	Not Hispanic or Latino	White
2	2	wP	Female	Not Hispanic or Latino	White
3	3	wP	Female	Unknown	White
4	4	wP	Male	Not Hispanic or Latino	Asian
5	5	wP	Male	Not Hispanic or Latino	Asian
6	6	wP	Female	Not Hispanic or Latino	White

	year_of_birth	date_of_boost	dataset	age	age_at_boost
1	1986-01-01	2016-09-12	2020_dataset	40.40246	30.69678
2	1968-01-01	2019-01-28	2020_dataset	58.40383	51.07461
3	1983-01-01	2016-10-10	2020_dataset	43.40315	33.77413
4	1988-01-01	2016-08-29	2020_dataset	38.40383	28.65982
5	1991-01-01	2016-08-29	2020_dataset	35.40315	25.65914
6	1988-01-01	2016-10-10	2020_dataset	38.40383	28.77481

We'd use `full_join` here since there are multiple specimens for each subject so you need to preserve that data.

```
meta <- full_join(specimen, subject)
```

Joining with ``by = join_by(subject_id)``

```
dim(meta)
```

```
[1] 1504  15
```

```
head(meta)
```

```
specimen_id subject_id actual_day_relative_to_boost
1           1           1                        -3
2           2           1                         1
3           3           1                         3
4           4           1                         7
5           5           1                        11
6           6           1                        32
planned_day_relative_to_boost specimen_type visit infancy_vac biological_sex
1                             0         Blood      1         wP         Female
2                             1         Blood      2         wP         Female
3                             3         Blood      3         wP         Female
4                             7         Blood      4         wP         Female
5                            14         Blood      5         wP         Female
6                            30         Blood      6         wP         Female
ethnicity race year_of_birth date_of_boost dataset
1 Not Hispanic or Latino White 1986-01-01 2016-09-12 2020_dataset
2 Not Hispanic or Latino White 1986-01-01 2016-09-12 2020_dataset
3 Not Hispanic or Latino White 1986-01-01 2016-09-12 2020_dataset
4 Not Hispanic or Latino White 1986-01-01 2016-09-12 2020_dataset
5 Not Hispanic or Latino White 1986-01-01 2016-09-12 2020_dataset
6 Not Hispanic or Latino White 1986-01-01 2016-09-12 2020_dataset
age age_at_boost
1 40.40246 30.69678
2 40.40246 30.69678
3 40.40246 30.69678
4 40.40246 30.69678
5 40.40246 30.69678
6 40.40246 30.69678
```

Q10. Now using the same procedure join meta with titer data so we can further analyze this data in terms of time of visit aP/wP, male/female etc.

```
abdata <- inner_join(meta, titer)
```

```
Joining with `by = join_by(specimen_id)`
```

```
dim(abdata)
```

```
[1] 52576    22
```

Q11. How many specimens (i.e. entries in abdata) do we have for each isotype?

```
table(abdata$isotype)
```

```
  IgE  IgG IgG1 IgG2 IgG3 IgG4
6698 5389 10117 10124 10124 10124
```

Shown above for each antibody isotype.

Q12. What are the different \$dataset values in abdata and what do you notice about the number of rows for the most “recent” dataset?

```
table(abdata$dataset)
```

```
2020_dataset 2021_dataset 2022_dataset 2023_dataset
        31520         8085         7301         5670
```

Looks like the most recent datasets have lots lower number of rows compared to the 2020 dataset, which has the most amount.

Examining IgG Ab titer levels

```
igg <- abdata %>% filter(isotype == "IgG")
head(igg)
```

```
  specimen_id subject_id actual_day_relative_to_boost
1           1           1                        -3
2           1           1                        -3
3           1           1                        -3
4           2           1                         1
5           2           1                         1
6           2           1                         1
```

	planned_day_relative_to_boost	specimen_type	visit	infancy_vac	biological_sex
1	0	Blood	1	wP	Female
2	0	Blood	1	wP	Female
3	0	Blood	1	wP	Female
4	1	Blood	2	wP	Female
5	1	Blood	2	wP	Female
6	1	Blood	2	wP	Female

	ethnicity	race	year_of_birth	date_of_boost	dataset
1	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset
2	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset
3	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset
4	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset
5	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset
6	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset


	age	age_at_boost	isotype	is_antigen_specific	antigen	MFI
1	40.40246	30.69678	IgG	TRUE	PT	68.56614
2	40.40246	30.69678	IgG	TRUE	PRN	332.12718
3	40.40246	30.69678	IgG	TRUE	FHA	1887.12263
4	40.40246	30.69678	IgG	TRUE	PT	41.38442
5	40.40246	30.69678	IgG	TRUE	PRN	174.89761
6	40.40246	30.69678	IgG	TRUE	FHA	246.00957

	MFI_normalised	unit	lower_limit_of_detection
1	3.736992	IU/ML	0.530000
2	2.602350	IU/ML	6.205949
3	34.050956	IU/ML	4.679535
4	2.255534	IU/ML	0.530000
5	1.370393	IU/ML	6.205949
6	4.438960	IU/ML	4.679535

Q13. Complete the following code to make a summary boxplot of Ab titer levels (MFI) for all antigens:

```
ggplot(igg) +
  aes(MFI_normalised, antigen) +
  geom_boxplot() +
  xlim(0,75) +
  facet_wrap(vars(igg$visit), nrow=2)
```

Warning: Removed 5 rows containing non-finite outside the scale range (`stat\_boxplot()`).



Q14. What antigens show differences in the level of IgG antibody titers recognizing them over time? Why these and not others?


It looks like FIM2/3 has a clear increase, alongside PT to some small extent. FHA shows a small increase but it diminishes over visit time.

FIM2 and 3 are both surface proteins on pertussis, therefore it makes sense for antibodies to attach it. FHA is filamentous hemagglutinin, surface protein used for binding to hosts. Finally PT is the toxin complex for pertussis therefore it's useful to target this as well since it's outside the cell. TT is tetanus toxin which is unrelated, DT is diphtheria which is unrelated to pertussis as well, PRN is pertactin autotransporter, which is why it shows a slight increase but not too much. OVA is ovalbumin, which is an internal component of many cells.

Q15. Filter to pull out only two specific antigens for analysis and create a boxplot for each. You can chose any you like. Below I picked a “control” antigen (“OVA”, that is not in our vaccines) and a clear antigen of interest (“PT”, Pertussis Toxin, one of the key virulence factors produced by the bacterium *B. pertussis*).


```
filter(igg, antigen=="OVA") %>%
  ggplot() +
  aes(MFI_normalised, col=infancy_vac) +
  geom_boxplot(show.legend = TRUE) +
  facet_wrap(vars(visit)) +
  theme_bw() +
```

```
labs(title="OVA Antigen IgG Levels by Vaccine Type and Visit",
      x="MFI Normalised")
```




```
filter(igg, antigen=="PT") %>%
  ggplot() +
  aes(MFI_normalised, col=infancy_vac) +
  geom_boxplot(show.legend = TRUE) +
  facet_wrap(vars(visit)) +
  theme_bw() +
  labs(title="PT Antigen IgG Levels by Vaccine Type and Visit",
        x="MFI Normalised")
```

PT Antigen IgG Levels by Vaccine Type and Visit



```
filter(igg, antigen=="FIM2/3") %>%
  ggplot() +
  aes(MFI_normalised, col=infancy_vac) +
  geom_boxplot(show.legend = TRUE) +
  facet_wrap(vars(visit)) +
  theme_bw() +
  labs(title="FIM2/3 Antigen IgG Levels by Vaccine Type and Visit",
        x="MFI Normalised")
```

FIM2/3 Antigen IgG Levels by Vaccine Type and Visit



Q16. What do you notice about these two antigens time courses and the PT data in particular?

It seems like for PT, aP levels are a bit lower in terms of its median compared to wP, while still overlapping a great deal. FIM2/3 on the other hand doesn't show such a great difference as PT does.

Q17. Do you see any clear difference in aP vs. wP responses?

It seemed like for PT, wP is faster to act in many individuals and is much bigger of an antibody response. In FIM2/3 this seems to be somewhat true however its response dies a bit faster than the aP response.

Next is the time course for 2021 data only

```
abdata.21 <- abdata %>% filter(dataset == "2021_dataset")

abdata.21 %>%
  filter(isotype == "IgG", antigen == "PT") %>%
  ggplot() +
    aes(x=planned_day_relative_to_boost,
         y=MFI_normalised,
         col=infancy_vac,
         group=subject_id) +
```




```

geom_point() +
geom_line() +
geom_vline(xintercept=0, linetype="dashed") +
geom_vline(xintercept=14, linetype="dashed") +
labs(title="2021 dataset IgG PT",
      subtitle = "Dashed lines indicate day 0 (pre-boost) and 14 (apparent peak levels)")

```

2021 dataset IgG PT

Dashed lines indicate day 0 (pre-boost) and 14 (apparent peak levels)



Q18. Does this trend look similar for the 2020 dataset?

```

abdata.20 <- abdata %>% filter(dataset == "2020_dataset")


abdata.20 %>%
  filter(isotype == "IgG", antigen == "PT") %>%
  ggplot() +
    aes(x=planned_day_relative_to_boost,
         y=MFI_normalised,
         col=infancy_vac,
         group=subject_id) +
  geom_point() +
  geom_line() +
  geom_vline(xintercept=0, linetype="dashed") +
  geom_vline(xintercept=14, linetype="dashed") +

```

```
labs(title="2020 dataset IgG PT",
      subtitle = "Dashed lines indicate day 0 (pre-boost) and 14 (apparent peak levels)") +
xlim(-10, 125)
```

Warning: Removed 3 rows containing missing values or values outside the scale range (``geom_point()``).

Warning: Removed 3 rows containing missing values or values outside the scale range (``geom_line()``).



It does seem similar but it seems more complex, since in the 2020 dataset there are more datapoints there is a lot of variation. It does seem that overall the MFI-normalized drops but there is no advantage in either aP or wP, or rather it would be impossible to differentiate without any sort of statistical comparison.

Obtaining CMI-PB RNASeq data

```
url <- "https://www.cmi-pb.org/api/v2/rnaseq?versioned_ensembl_gene_id=eq.ENSNG00000211896.7"
rna <- read_json(url, simplifyVector = TRUE)
```

```
meta <- inner_join(specimen, subject)
```


Joining with `by = join\_by(subject\_id)`

```
ssrna <- inner_join(meta, rna)
```

Joining with `by = join\_by(specimen\_id)`

Q19. Make a plot of the time course of gene expression for IGHG1 gene (i.e. a plot of visit vs. tpm).

```
ggplot(ssrna) +  
  aes(x=visit, y=tpm, group=subject_id) +  
  geom_point() +  
  geom_line(alpha=0.2)
```



Q20.: What do you notice about the expression of this gene (i.e. when is it at it's maximum level)?


Maximum level seems to peak at visit 4, which seems to match up with the idea that IgG expression is more of a later phenomena whereas IgM is more common early on. It takes time for recombination to make IgG.

Q21. Does this pattern in time match the trend of antibody titer data? If not, why not?

This does seem to match the prior box plots as in those the IgG count usually grew for visit 4, which matches the expression found above.

Doing box plots:

```
ggplot(ssrna) +  
  aes(tpm, col=infancy_vac) +  
  geom_boxplot() +  
  facet_wrap(vars(visit))
```



Also interesting, visit 4 has a lot of expression of aP and a slight increase in wP expression.

Q. Is RNA-Seq expression levels predictive of Ab titers?

It does seem so as increase in expression seems to match up with the increase in titers.

Q. What differentiates aP vs. wP primed individuals?

It seems like aP primed individuals have somewhat higher RNA expression of IGHG1 alongside worse antibody match to PT and slightly better for FIM2/3. This is very hard to differentiate however.

Q. What about decades after their first immunization? Do you know?

The main difference is that wP individuals are older than aP, but we weren't able to look at long term data it seems on RNA seq or antibody titers or how long lasting they are.